

# Stack Overflow questions


JOINING DATA WITH DPLYR




**Chris Cardillo**  
Data Scientist at DataCamp

# The joining verbs


Inner join




Left join




Right join




Full join



Semi join




Anti join



# Can dplyr join on multiple columns or composite key?


Asked 4 years, 9 months ago   Active 1 year ago   Viewed 93k times

  
89

I realize that `dplyr` v3.0 allows you to join on different variables:

```
left_join(x, y, by = c("a" = "b"))
```

will match `x.a` to `y.b`

  
25

However, is it possible to join on a combination of variables or do I have to add a composite key beforehand?

Something like this:


```
left_join(x, y, by = c("a c" = "b d"))
```


to match the concatenation of `[ x.a and x.c ]` to `[ y.b and y.d ]`

r

dplyr

share edit close flag


edited Jul 18 '18 at 15:16  
 MusTheDataGuy  
2,462 ● 18 ● 61 ● 98



asked Oct 28 '14 at 15:07  
 JasonAizkalns  
13.3k ● 4 ● 36 ● 87

[add a comment](#)

1 Answer

active   oldest   votes

  
154

You can pass a named vector of length greater than 1 to the `by` argument of `left_join()`:

```
library(dplyr)

d1 <- data_frame(
  x = letters[1:3],
  y = LETTERS[1:3],
  a = rnorm(3)
)
```

# The questions table

questions

```
# A tibble: 294,735 x 3
      id creation_date score
  <int> <date>         <int>
1 22557677 2014-03-21         1
2 22557707 2014-03-21         2
3 22558084 2014-03-21         2
4 22558395 2014-03-21         2
5 22558613 2014-03-21         0
6 22558677 2014-03-21         2
7 22558887 2014-03-21         8
8 22559180 2014-03-21         1
9 22559312 2014-03-21         0
10 22559322 2014-03-21         2
# ... with 294,725 more rows
```

# The question\_tags and tags tables

question\_tags

```
# A tibble: 497,153 x 2
  question_id tag_id
      <int>   <int>
1    22557677     18
2    22557677    139
3    22557677  16088
4    22557677   1672
5    22558084   6419
6    22558084  92764
7    22558395   5569
8    22558395    134
9    22558395   9412
10   22558395  18621
# ... with 497,143 more rows
```

tags

```
# A tibble: 48,299 x 2
      id tag_name
    <dbl> <chr>
1  124399 laravel-dusk
2  124402 spring-cloud-vault-config
3  124404 spring-vault
4  124405 apache-bahir
5  124407 astc
6  124408 simulacrum
7  124410 angularartics2
8  124411 django-rest-viewssets
9  124414 react-native-lightbox
10 124417 java-module
# ... with 48,289 more rows
```

# Joining question\_tags with questions

```
questions %>%  
  inner_join(question_tags, by = c("id" = "question_id"))
```

# Joining tags

```
questions_with_tags <- questions %>%  
  inner_join(question_tags, by = c("id" = "question_id")) %>%  
  inner_join(tags, by = c("tag_id" = "id"))
```

```
questions_with_tags
```

```
# A tibble: 497,153 x 5  
      id creation_date score tag_id tag_name  
  <int> <date>         <int> <dbl> <chr>  
1 22557677 2014-03-21      1     18 regex  
2 22557677 2014-03-21      1    139 string  
3 22557677 2014-03-21      1 16088 time-complexity  
4 22557677 2014-03-21      1   1672 backreference  
5 22558084 2014-03-21      2   6419 time-series  
6 22558084 2014-03-21      2  92764 panel-data  
7 22558395 2014-03-21      2   5569 function  
8 22558395 2014-03-21      2    134 sorting  
9 22558395 2014-03-21      2   9412 vectorization  
10 22558395 2014-03-21      2 18621 operator-precedence  
# ... with 497,143 more rows
```

# Most common tags

```
questions_with_tags %>%  
  count(tag_name, sort = TRUE)
```

```
# A tibble: 7,840 x 2  
  tag_name      n  
  <chr>      <int>  
1 ggplot2    28228  
2 dataframe  18874  
3 shiny      14219  
4 dplyr       14039  
5 plot       11315  
6 data.table   8809  
7 matrix      6205  
8 loops       5149  
9 regex       4912  
10 function   4892  
# ... with 7,830 more rows
```



# Joining questions and answers

JOINING DATA WITH DPLYR



**Chris Cardillo**  
Data Scientist at DataCamp

# The answers table

answers

```
# A tibble: 380,643 x 4
      id creation_date question_id score
  <int> <date>          <int> <int>
1 39143713 2016-08-25      39143518     3
2 39143869 2016-08-25      39143518     1
3 39143935 2016-08-25      39142481     0
4 39144014 2016-08-25      39024390     0
5 39144252 2016-08-25      39096741     6
6 39144375 2016-08-25      39143885     5
7 39144430 2016-08-25      39144077     0
8 39144625 2016-08-25      39142728     1
9 39144794 2016-08-25      39043648     0
10 39145033 2016-08-25      39133170     1
# ... with 380,633 more rows
```


# The question ID

```
questions %>%  
  inner_join(answers, by = c("id" = "question_id"))
```


```
# A tibble: 380,643 x 6  
      id creation_date.x score.x      id.y creation_date.y score.y  
  <int> <date>          <int>   <int> <date>          <int>  
1 22557677 2014-03-21          1 22560670 2014-03-21          2  
2 22557707 2014-03-21          2 22558516 2014-03-21          1  
3 22557707 2014-03-21          2 22558726 2014-03-21          4  
4 22558084 2014-03-21          2 22558085 2014-03-21          0  
5 22558084 2014-03-21          2 22606545 2014-03-24          1  
6 22558084 2014-03-21          2 22610396 2014-03-24          5  
7 22558084 2014-03-21          2 34374729 2015-12-19          0  
8 22558395 2014-03-21          2 22559327 2014-03-21          1  
9 22558395 2014-03-21          2 22560102 2014-03-21          2  
10 22558395 2014-03-21          2 22560288 2014-03-21          2  
# ... with 380,633 more rows
```

# The joining verbs


Inner join




Left join




Right join




Full join



Semi join



Anti join



# The `bind_rows` verb

JOINING DATA WITH DPLYR



**Chris Cardillo**

Data Scientist at DataCamp

# Comparing tables

## questions

```
# A tibble: 294,735 x 3
      id creation_date score
  <int> <date>         <int>
1 22557677 2014-03-21         1
2 22557707 2014-03-21         2
3 22558084 2014-03-21         2
4 22558395 2014-03-21         2
5 22558613 2014-03-21         0
6 22558677 2014-03-21         2
7 22558887 2014-03-21         8
8 22559180 2014-03-21         1
9 22559312 2014-03-21         0
10 22559322 2014-03-21         2
# ... with 294,725 more rows
```

## answers

```
# A tibble: 380,635 x 4
      id creation_date question_id score
  <int> <date>         <int> <int>
1 39143713 2016-08-25      39143518     3
2 39143869 2016-08-25      39143518     1
3 39143935 2016-08-25      39142481     0
4 39144014 2016-08-25      39024390     0
5 39144252 2016-08-25      39096741     6
6 39144375 2016-08-25      39143885     5
7 39144430 2016-08-25      39144077     0
8 39144625 2016-08-25      39142728     1
9 39144794 2016-08-25      39043648     0
10 39145033 2016-08-25      39133170     1
# ... with 380,625 more rows
```

# Binding rows

```
questions %>%  
  bind_rows(answers)
```

```
# A tibble: 675,370 x 4  
      id creation_date score question_id  
  <int> <date>         <int>      <int>  
1 22557677 2014-03-21         1         NA  
2 22557707 2014-03-21         2         NA  
3 22558084 2014-03-21         2         NA  
4 22558395 2014-03-21         2         NA  
5 22558613 2014-03-21         0         NA  
6 22558677 2014-03-21         2         NA  
7 22558887 2014-03-21         8         NA  
8 22559180 2014-03-21         1         NA  
9 22559312 2014-03-21         0         NA  
10 22559322 2014-03-21         2         NA  
# ... with 675,360 more rows
```

# Using bind rows

```
questions_type <- questions %>%  
  mutate(type = "question")
```

```
answers_type <- answers %>%  
  mutate(type = "answer")
```

```
posts <- bind_rows(questions_type, answers_type)  
posts
```

```
# A tibble: 675,370 x 5  
      id creation_date score type      question_id  
  <int> <date>         <int> <chr>         <int>  
1 22557677 2014-03-21         1 question        NA  
2 22557707 2014-03-21         2 question        NA  
3 22558084 2014-03-21         2 question        NA  
4 22558395 2014-03-21         2 question        NA  
5 22558613 2014-03-21         0 question        NA  
6 22558677 2014-03-21         2 question        NA  
7 22558887 2014-03-21         8 question        NA  
8 22559180 2014-03-21         1 question        NA  
9 22559312 2014-03-21         0 question        NA  
10 22559322 2014-03-21         2 question        NA  
# ... with 675,360 more rows
```



# Aggregating

```
posts %>%  
  group_by(type) %>%  
  summarize(average_score = mean(score))
```

```
# A tibble: 2 x 2  
  type      average_score  
  <chr>         <dbl>  
1 answer         2.88  
2 question       1.90
```

# Creating date variable

```
library(lubridate)
```

```
posts %>%
```

```
  mutate(year = year(creation_date))
```

```
# A tibble: 675,370 x 6
```

	id	creation_date	score	type	question_id	year
	<int>	<date>	<int>	<chr>	<int>	<dbl>
1	22557677	2014-03-21	1	question	NA	2014
2	22557707	2014-03-21	2	question	NA	2014
3	22558084	2014-03-21	2	question	NA	2014
4	22558395	2014-03-21	2	question	NA	2014
5	22558613	2014-03-21	0	question	NA	2014
6	22558677	2014-03-21	2	question	NA	2014
7	22558887	2014-03-21	8	question	NA	2014
8	22559180	2014-03-21	1	question	NA	2014
9	22559312	2014-03-21	0	question	NA	2014
10	22559322	2014-03-21	2	question	NA	2014

```
# ... with 675,360 more rows
```

# Counting date variable

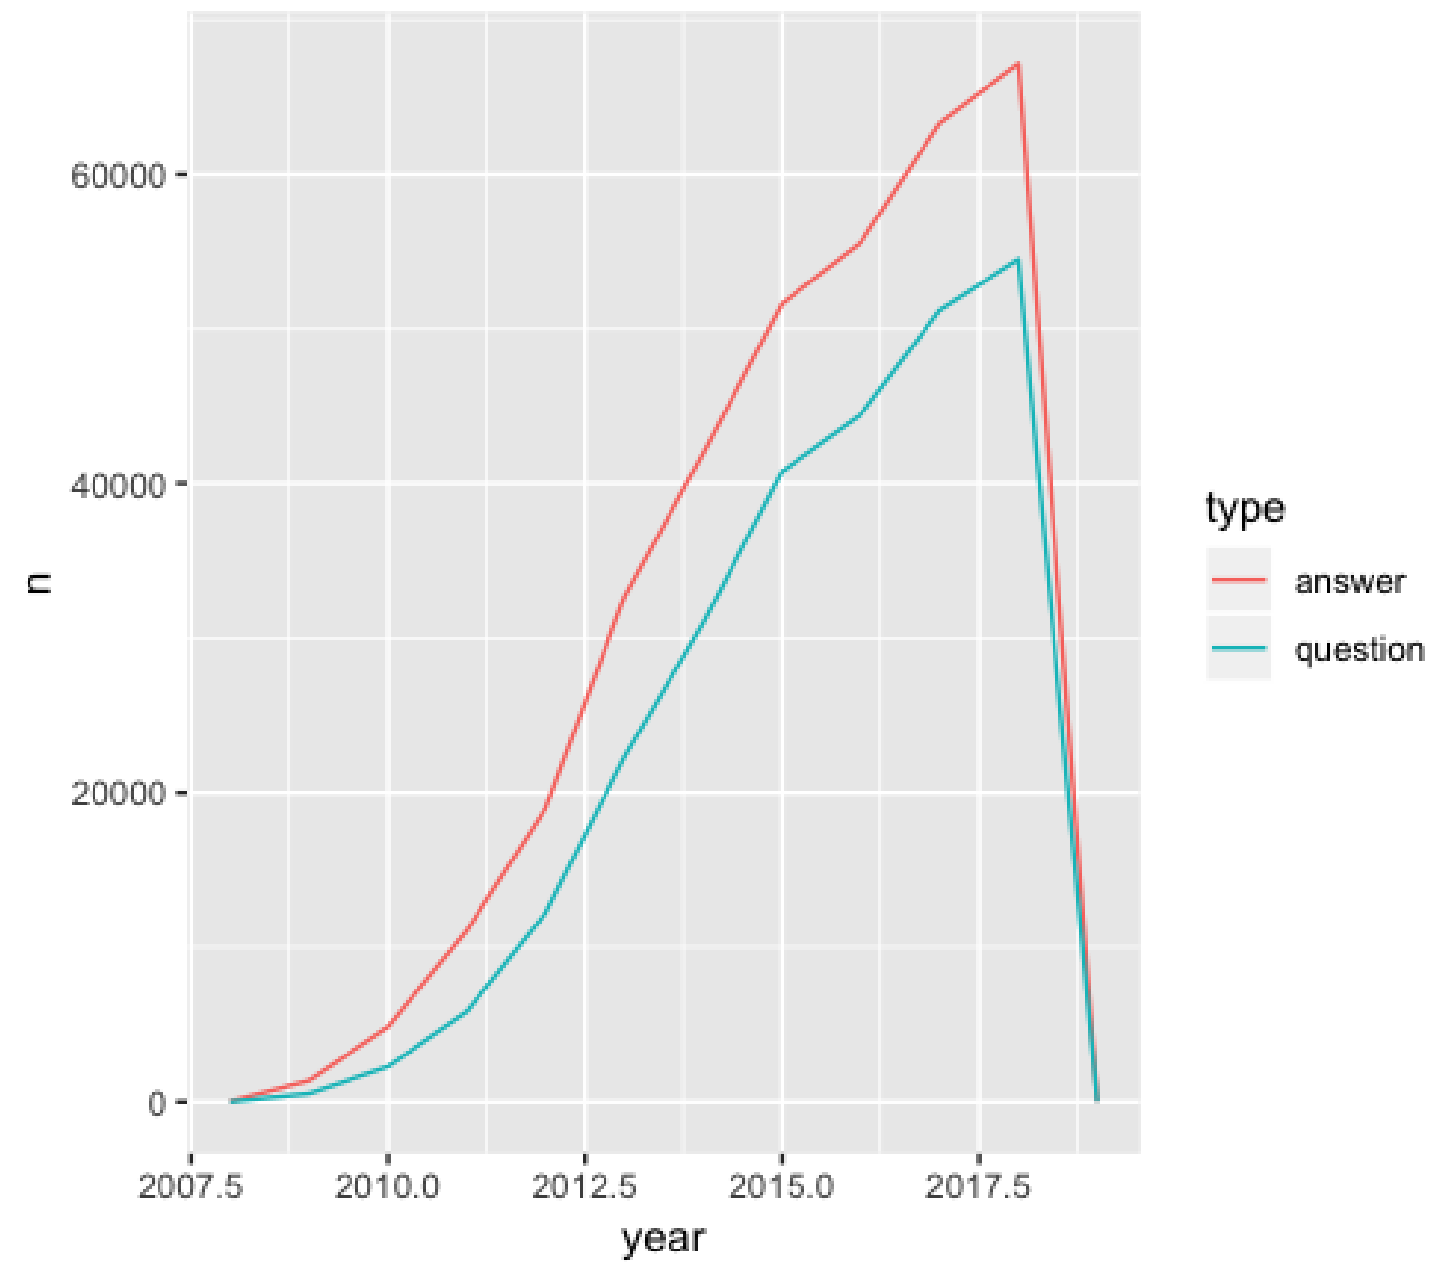
```
posts %>%  
  mutate(year = year(creation_date)) %>%  
  count(year, type)
```

```
# A tibble: 24 x 3  
   year type      n  
   <dbl> <chr>   <int>  
1  2008 answer     27  
2  2008 question    8  
3  2009 answer  1356  
4  2009 question  524  
5  2010 answer  4846  
6  2010 question 2264  
7  2011 answer 11077  
8  2011 question 5837  
9  2012 answer 18967  
10 2012 question 12210  
# ... with 14 more rows
```

# Plotting date variable


```
questions_answers_year <- posts %>%  
  mutate(year = year(creation_date)) %>%  
  count(year, type)  
  
ggplot(questions_answers_year, aes(year, n, color = type)) +  
  geom_line()
```

# The posts plot




# The joining verbs


Inner join




Left join




Right join




Full join



Semi join




Anti join




# The mutating joins

## Inner join




# The mutating joins

## Inner join




## Left join






# The mutating joins


## Inner join



## Left join




## Right join




# The mutating joins


## Inner join




## Left join



## Full join



## Right join




# The filtering joins

Semi join


first table

second table




# The filtering joins

Semi join



Anti join



# Congratulations!

